CoderUJB · frozen 100 tasks · A/B reviewed results

FCG

	A	B	C	D	E
CodeLlama-13B	2/20 (10%) 1 unknown	3/20 (15%) Δ +5.3 pp (19 pairs)	Not run	Not run	Not run
Qwen2.5-Coder-32B	2/20 (10%)	3/20 (15%) Δ +5.0 pp	Not run	Not run	Not run
gpt-oss-120b	3/20 (15%)	5/20 (25%) Δ +10.0 pp	Not run	Not run	Not run
GPT-6.1 Sol	14/20 (70%)	15/20 (75%) Δ +5.0 pp	Not run	Not run	Not run

CTG

	A	B	C	D	E
CodeLlama-13B	3/20 (15%)	5/20 (25%) Δ +10.0 pp	Not run	Not run	Not run
Qwen2.5-Coder-32B	6/20 (30%)	3/20 (15%) Δ -15.0 pp	Not run	Not run	Not run
gpt-oss-120b	5/20 (25%)	7/20 (35%) Δ +10.0 pp	Not run	Not run	Not run
GPT-6.1 Sol	16/20 (80%)	13/20 (65%) Δ -15.0 pp	Not run	Not run	Not run

ITG

	A	B	C	D	E
CodeLlama-13B	0/20 (0%)	0/20 (0%) Δ +0.0 pp	Not run	Not run	Not run
Qwen2.5-Coder-32B	1/20 (5%)	3/20 (15%) Δ +10.0 pp	Not run	Not run	Not run
gpt-oss-120b	1/20 (5%)	2/20 (10%) Δ +5.0 pp	Not run	Not run	Not run
GPT-6.1 Sol	10/20 (50%)	10/20 (50%) Δ +0.0 pp	Not run	Not run	Not run

APR

	A	B	C	D	E
CodeLlama-13B	0/20 (0%)	2/20 (10%) Δ +10.0 pp	Not run	Not run	Not run
Qwen2.5-Coder-32B	2/20 (10%) 1 unknown	3/20 (15%) Δ +5.3 pp (19 pairs)	Not run	Not run	Not run
gpt-oss-120b	6/20 (30%)	6/20 (30%) Δ +0.0 pp	Not run	Not run	Not run
GPT-6.1 Sol	14/20 (70%)	16/20 (80%) Δ +10.0 pp	Not run	Not run	Not run

DD

	A	B	C	D	E
CodeLlama-13B	10/20 (50%)	9/20 (45%) Δ -5.0 pp	Not run	Not run	Not run
Qwen2.5-Coder-32B	9/20 (45%)	10/20 (50%) Δ +5.0 pp	Not run	Not run	Not run
gpt-oss-120b	12/20 (60%)	10/20 (50%) Δ -10.0 pp	Not run	Not run	Not run
GPT-6.1 Sol	16/20 (80%)	16/20 (80%) Δ +0.0 pp	Not run	Not run	Not run

Reviewed v2 · 800 saved formal answers
798 determinate answers · 398 determinate A/B pairs

A: official task prompt
B: fixed output constraint
C–E: NOT RUN (gray)

Cells: passes / 20 planned; unknowns remain unknown.
Color uses the observed lower bound, not known-only accuracy.
B delta: paired B − A; each task has 19 or 20 known pairs.

CTG correction: CodeLlama B 4 → 5 passes.
The added pass is an empty test credited for enum initialization.
Benchmark pass does not establish test usefulness.
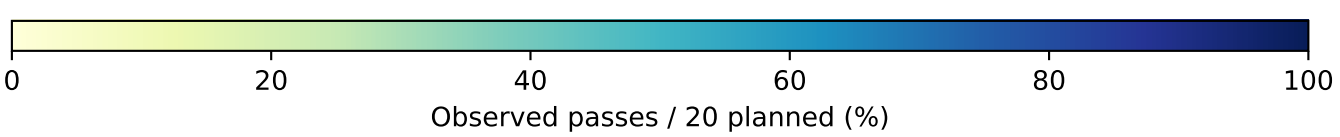
Exploratory follow-up after seeing A;
one generation per cell.
Different model formats and output budgets; not a causal estimate of model age, nor the full paper reproduction.